

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
data = pd.read_csv("/content/hotel_bookings.csv")
data.head()
```

```
missing_values = data.isnull().sum()
print("Missing Value: ")
print(missing_values)
```

```
data.dropna(inplace=True)
```

```
duplicate_rows = data.duplicated().sum()
print("Duplicated rows: ", duplicate_rows)
data.drop_duplicates(inplace=True)
```

```
Missing Value:
hotel                                0
is_canceled                          0
lead_time                            0
arrival_date_year                     0
arrival_date_month                    0
arrival_date_week_number              0
arrival_date_day_of_month              0
stays_in_weekend_nights                0
stays_in_week_nights                  0
adults                                0
children                              4
babies                                0
meal                                  0
country                             488
market_segment                        0
distribution_channel                  0
is_repeated_guest                     0
previous_cancellations                 0
previous_bookings_not_canceled         0
reserved_room_type                    0
assigned_room_type                     0
booking_changes                       0
deposit_type                           0
agent                               16340
company                             112593
days_in_waiting_list                  0
customer_type                          0
adr                                    0
required_car_parking_spaces            0
total_of_special_requests               0
reservation_status                     0
reservation_status_date                 0
dtype: int64
Duplicated rows: 28
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 189 entries, 2392 to 116451
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype

```

```

---  -----
0  hotel                189 non-null    object
1  is_canceled          189 non-null    int64
2  lead_time            189 non-null    int64
3  arrival_date_year    189 non-null    int64
4  arrival_date_month   189 non-null    object
5  arrival_date_week_number 189 non-null    int64
6  arrival_date_day_of_month 189 non-null    int64
7  stays_in_weekend_nights 189 non-null    int64
8  stays_in_week_nights  189 non-null    int64
9  adults               189 non-null    int64
10 children             189 non-null    float64
11 babies               189 non-null    int64
12 meal                 189 non-null    object
13 country              189 non-null    object
14 market_segment       189 non-null    object
15 distribution_channel  189 non-null    object
16 is_repeated_guest    189 non-null    int64
17 previous_cancellations 189 non-null    int64
18 previous_bookings_not_canceled 189 non-null    int64
19 reserved_room_type   189 non-null    object
20 assigned_room_type    189 non-null    object
21 booking_changes       189 non-null    int64
22 deposit_type          189 non-null    object
23 agent                189 non-null    float64
24 company              189 non-null    float64
25 days_in_waiting_list  189 non-null    int64
26 customer_type         189 non-null    object
27 adr                  189 non-null    float64
28 required_car_parking_spaces 189 non-null    int64
29 total_of_special_requests 189 non-null    int64
30 reservation_status    189 non-null    object
31 reservation_status_date 189 non-null    object
dtypes: float64(4), int64(16), object(12)
memory usage: 48.7+ KB

```

```
data.nunique()
```

```

hotel                2
is_canceled          2
lead_time            50
arrival_date_year     3
arrival_date_month    12
arrival_date_week_number 38
arrival_date_day_of_month 29
stays_in_weekend_nights 8
stays_in_week_nights  13
adults               3
children             3
babies               1
meal                 4
country              18
market_segment        6
distribution_channel  4
is_repeated_guest     2
previous_cancellations 2
previous_bookings_not_canceled 6
reserved_room_type    5
assigned_room_type     9
booking_changes        7
deposit_type          3
agent                28

```

company	46
days_in_waiting_list	1
customer_type	4
adr	70
required_car_parking_spaces	2
total_of_special_requests	4
reservation_status	3
reservation_status_date	76
dtype: int64	

data.dtypes

hotel	object
is_canceled	int64
lead_time	int64
arrival_date_year	int64
arrival_date_month	object
arrival_date_week_number	int64
arrival_date_day_of_month	int64
stays_in_weekend_nights	int64
stays_in_week_nights	int64
adults	int64
children	float64
babies	int64
meal	object
country	object
market_segment	object
distribution_channel	object
is_repeated_guest	int64
previous_cancellations	int64
previous_bookings_not_canceled	int64
reserved_room_type	object
assigned_room_type	object
booking_changes	int64
deposit_type	object
agent	float64
company	float64
days_in_waiting_list	int64
customer_type	object
adr	float64
required_car_parking_spaces	int64
total_of_special_requests	int64
reservation_status	object
reservation_status_date	object
dtype: object	

data.describe()

	is_canceled	lead_time	arrival_date_year	arrival_date_week_number	arrival_date
count	189.000000	189.000000	189.000000	189.000000	189.000000
mean	0.084656	33.883598	2015.486772	37.555556	37.555556
std	0.279109	50.797022	0.740963	13.317854	13.317854
min	0.000000	0.000000	2015.000000	1.000000	1.000000
25%	0.000000	11.000000	2015.000000	33.000000	33.000000
50%	0.000000	24.000000	2015.000000	45.000000	45.000000
75%	0.000000	36.000000	2016.000000	46.000000	46.000000
max	1.000000	364.000000	2017.000000	53.000000	53.000000

```

to_drop = ['country', 'market_segment',
           'distribution_channel', 'reserved_room_type', 'assigned_room_type',
           'deposit_type', 'customer_type', 'reservation_status',
           'reservation_status_date', 'hotel']
X=data.drop(to_drop,axis=1)
X.head()

```

	is_canceled	lead_time	arrival_date_year	arrival_date_month	arrival_date
2392	0	6	2015	October	2015-10-06
2697	0	24	2015	October	2015-10-24
2867	0	24	2015	November	2015-11-24
2877	0	24	2015	November	2015-11-24
2878	0	24	2015	November	2015-11-24

5 rows x 22 columns

```

cat_cols=X.columns[X.dtypes == "object"]
num_cols=X.columns[(X.dtypes == "float64")|(X.dtypes == "int64")]

```

```
num_cols
```

```

Index(['is_canceled', 'lead_time', 'arrival_date_year',
       'arrival_date_week_number', 'arrival_date_day_of_month',
       'stays_in_weekend_nights', 'stays_in_week_nights', 'adults', 'children',
       'babies', 'is_repeated_guest', 'previous_cancellations',
       'previous_bookings_not_canceled', 'booking_changes', 'agent', 'company',
       'days_in_waiting_list', 'adr', 'required_car_parking_spaces',
       'total_of_special_requests'],
      dtype='object')

```

```
cat_cols
```

```

Index(['arrival_date_month', 'meal'], dtype='object')

```

```

from sklearn import model_selection
y = data['is_canceled']
X_train,X_test,y_train,y_test = model_selection.train_test_split(X,y,test_size=0.25,stratify =
print("train data has"+ str(X_train.shape[0])+"observation with"+ str(X_train.shape[1])+ "feat
print("test data has"+str(X_test.shape[0])+"observation with"+str(X_test.shape[1])+ "features'

```

```

train data has141observation with22features
test data has48observation with22features

```

```
data_train.head()
```

	hotel	is_canceled	lead_time	arrival_date_year	arrival_date_month	arrival_
0	Resort Hotel	-0.304997	0.056076	-0.699663	November	
1	Resort Hotel	-0.304997	4.598244	2.003144	June	
2	City Hotel	-0.304997	-0.712968	-0.699663	August	
3	Resort Hotel	-0.304997	-0.592805	2.003144	January	
4	Resort Hotel	-0.304997	-0.688935	0.651741	July	

5 rows x 36 columns

```

from sklearn.preprocessing import OneHotEncoder
def OneHotEncoding(df,enc,categories):
    transformed = pd.DataFrame(enc.transform(df[categories]).toarray(),columns=enc.get_feature_r
    return pd.concat([df.reset_index(drop=True),transformed],axis=1).drop(categories,axis=1)

```

```

categories = ["arrival_date_month"]
enc_ohe = OneHotEncoder()
enc_ohe.fit(X_train[categories])

```

```

X_train=OneHotEncoding(X_train,enc_ohe,categories)
X_test=OneHotEncoding(X_test,enc_ohe,categories)

```

```

from sklearn.preprocessing import OrdinalEncoder
categories_order = [['FB', 'HB', 'BB','SC']]
enc_oe = OrdinalEncoder(categories=categories_order)
enc_oe.fit(X_train[["meal"]])

```

```

X_train["meal_encoded"] = enc_oe.transform(X_train[["meal"]])
X_test["meal_encoded"] = enc_oe.transform(X_test[["meal"]])

```

```

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(data_train[num_cols])
data_train[num_cols] = scaler.transform(data_train[num_cols])
data_test[num_cols] = scaler.transform(data_test[num_cols])

from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
classifier_logistic = LogisticRegression()
classifier_KNN = KNeighborsClassifier()
classifier_RF = RandomForestClassifier()

data_train_encoded = pd.get_dummies(data_train, columns=["arrival_date_month", "meal"])
data_test_encoded = pd.get_dummies(data_test, columns=["arrival_date_month", "meal"])
encoded_columns = data_train_encoded.columns
print(encoded_columns)

Index(['hotel', 'is_canceled', 'lead_time', 'arrival_date_year',
      'arrival_date_week_number', 'arrival_date_day_of_month',
      'stays_in_weekend_nights', 'stays_in_week_nights', 'adults', 'children',
      'babies', 'country', 'distribution_channel', 'is_repeated_guest',
      'previous_cancellations', 'previous_bookings_not_canceled',
      'reserved_room_type', 'assigned_room_type', 'booking_changes',
      'deposit_type', 'agent', 'company', 'days_in_waiting_list',
      'customer_type', 'adr', 'required_car_parking_spaces',
      'total_of_special_requests', 'reservation_status',
      'reservation_status_date', 'market_segment_Complementary',
      'market_segment_Corporate', 'market_segment_Direct',
      'market_segment_Offline TA/TO', 'market_segment_Online TA',
      'arrival_date_month_April', 'arrival_date_month_August',
      'arrival_date_month_December', 'arrival_date_month_February',
      'arrival_date_month_January', 'arrival_date_month_July',
      'arrival_date_month_June', 'arrival_date_month_March',
      'arrival_date_month_May', 'arrival_date_month_November',
      'arrival_date_month_October', 'arrival_date_month_September', 'meal_BB',
      'meal_HB', 'meal_SC'],
      dtype='object')

from sklearn.linear_model import LogisticRegression
classifier_logistic = LogisticRegression()
classifier_logistic.fit(X_train_processed, y_train)

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: (
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

▼ LogisticRegression

LogisticRegression()

```
X_test_processed = X_test.drop(columns=['meal'])
predictions = classifier_logistic.predict(X_test_processed)
```


```
X_test_processed = X_test.drop(columns=['meal'])
score = classifier_logistic.score(X_test_processed, y_test)
print("Accuracy on test set:", score)
```

Accuracy on test set: 0.8541666666666666

```
from sklearn.model_selection import GridSearchCV
def print_grid_search_metrics(gs):
    print("Best score:"+str(gs.best_score_))
    print("Best parameters set:")
    best_parameters = gs.best_params_
    for param_name in sorted(best_parameters.keys()):
        print(param_name+"-"+str(best_parameters[param_name]))
```

```
parameters = {
    'penalty': ('l2', 'l1'),
    'C': (0.01, 0.05, 0.1, 0.2, 1)
}
Grid_LR = GridSearchCV(LogisticRegression(solver='liblinear'), parameters, cv=5)
Grid_LR.fit(X_test_processed, y_test)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: l
warnings.warn(
```



```
print_grid_search_metrics(Grid_LR)
```


```
Best score:0.9177777777777777
Best parameters set:
C:0.01
penalty:l2
```

```
best_LR_model = Grid_LR.best_estimator_
```


```
best_LR_model.score(X_test_processed, y_test)
```

0.9166666666666666

```
LR_models = pd.DataFrame(Grid_LR.cv_results_)
res = (LR_models.pivot(index="param_penalty", columns="param_C", values="mean_test_score")
      )
_ = sns.heatmap(res, cmap="viridis", color="red")
```



```
parameters = {
    "n_neighbors": [1,3,5,7,9]
}
Grid_KNN = GridSearchCV(KNeighborsClassifier(),parameters,cv=5)
X_train_processed = X_train.drop(columns=['meal'])
Grid_KNN.fit(X_train_processed,y_train)
```




```
print_grid_search_metrics(Grid_KNN)
```

```
Best score:0.9150246305418721
Best parameters set:
n_neighbors:7
```

```
best_KNN_model=Grid_KNN.best_estimator_
```

```
parameters = {
    "n_estimators": [60,80,100],
    "max_depth": [1,5,10]
}
Grid_RF = GridSearchCV(RandomForestClassifier(),parameters,cv=5)
Grid_RF.fit(X_train_processed,y_train)
```

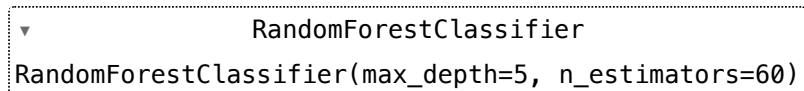



```
print_grid_search_metrics(Grid_RF)
```

```
Best score:0.9647783251231526
Best parameters set:
max_depth:5
n_estimators:60
```

```
best_RF_model = Grid_RF.best_estimator_
```

```
best_RF_model
```



```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```

```
def cal_evaluation(classifier,cm):
    tn=cm[0][0]
    fp=cm[0][1]
    fn=cm[1][0]
    tp=cm[1][1]
    accuracy = (tp+tn)/(tp+fp+fn+tn+0.0)
    precision = tp/(tp+fp+0.0)
    recall= tp/(tp+fn+0.0)
    print(classifier)
    print("Accuracy is" + str(accuracy))
    print("precision is" + str(precision))
    print("recall is " + str(recall))
    print()
```

```
def draw_confusion_matrices(confusion_matrices):
    class_names = ['Not','Churn']
    for cm in confusion_matrices:
        classifier, cm = cm[0], cm[1]
        cal_evaluation(classifier, cm)
```

```
confusion_matrices = [
    ("Random Forest", confusion_matrix(y_test,best_RF_model.predict(X_test_processed))),
    ("Logistic Regression", confusion_matrix(y_test,best_LR_model.predict(X_test_processed))),
    ("K nearest neighbor", confusion_matrix(y_test, best_KNN_model.predict(X_test_processed)))
]
```

```
draw_confusion_matrices(confusion_matrices)
```

```
Random Forest
Accuracy is0.9791666666666666
precision is1.0
recall is 0.75
```

```
Logistic Regression
Accuracy is0.9166666666666666
precision isnan
recall is 0.0
```

```
K nearest neighbor
Accuracy is0.9166666666666666
precision isnan
recall is 0.0
```

```
<ipython-input-328-ac5c1f3f75c2>:12: RuntimeWarning: invalid value encountered in scalar d
precision = tp/(tp+fp+0.0)
```

```
def draw_confusion_matrices(confusion_matrices):
    class_names = ["Not","Churn"]
    for cm in confusion_matrices:
        classifier, cm = cm[0],cm[1]
        cal_evaluation(classifier,cm)
```

```
from sklearn.metrics import roc_curve
from sklearn import metrics
```


```
y_pred_rf = best_RF_model.predict_proba(X_test_processed)[: , 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)
```

```
best_RF_model.predict_proba(X_test_processed)
```

```
array([[0.99583158, 0.00416842],
       [0.89324314, 0.10675686],
       [0.89944444, 0.10055556],
       [0.91198423, 0.08801577],
       [0.95608406, 0.04391594],
       [0.95608406, 0.04391594],
       [0.96303186, 0.03696814],
       [0.99764023, 0.00235977],
       [0.98088384, 0.01911616],
       [0.9890599 , 0.0109401 ],
       [0.97832574, 0.02167426],
       [0.96108197, 0.03891803],
       [0.98050505, 0.01949495],
       [0.98143764, 0.01856236],
       [0.9393484 , 0.0606516 ],
       [0.99155579, 0.00844421],
       [0.97608753, 0.02391247],
```

```
[0.92858362, 0.07141638],
[0.29974248, 0.70025752],
[0.82808141, 0.17191859],
[0.99027581, 0.00972419],
[0.45463654, 0.54536346],
[0.9133699 , 0.0866301 ],
[0.91104847, 0.08895153],
[0.99504294, 0.00495706],
[0.98799129, 0.01200871],
[0.90027343, 0.09972657],
[0.98013954, 0.01986046],
[0.92435516, 0.07564484],
[0.97970684, 0.02029316],
[0.95664744, 0.04335256],
[0.4126695 , 0.5873305 ],
[0.9268484 , 0.0731516 ],
[0.97534126, 0.02465874],
[0.96238517, 0.03761483],
[0.95646262, 0.04353738],
[0.50668129, 0.49331871],
[0.94534272, 0.05465728],
[0.99323185, 0.00676815],
[0.91743779, 0.08256221],
[0.93949134, 0.06050866],
[0.98106061, 0.01893939],
[0.97192086, 0.02807914],
[0.89324314, 0.10675686],
[0.90729167, 0.09270833],
[0.90506298, 0.09493702],
[0.99146873, 0.00853127],
[0.93401533, 0.06598467]])
```

```
import matplotlib.pyplot as plt
plt.figure(1)
plt.plot([0,1],[0,1],"k--")
plt.plot(fpr_rf, tpr_rf,label="RF")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.title("ROC curve = RF model")
plt.legend(loc = "best")
plt.show()
```



```
from sklearn import metrics
metrics.auc(fpr_rf,tpr_rf)
```

```
1.0
```

```
y_pred_lr=best_LR_model.predict_proba(X_test_processed)[: ,1]
fpr_lr,tpr_lr, thresh = roc_curve(y_test, y_pred_lr)
```

```
best_LR_model.predict_proba(X_test_processed)
```

```
array([[9.87239133e-01, 1.27608671e-02],
       [8.47691625e-01, 1.52308375e-01],
       [8.94441465e-01, 1.05558535e-01],
       [9.77802243e-01, 2.21977574e-02],
       [8.64704627e-01, 1.35295373e-01],
       [8.64704627e-01, 1.35295373e-01],
       [9.80547171e-01, 1.94528286e-02],
       [9.13940659e-01, 8.60593412e-02],
       [9.05397922e-01, 9.46020783e-02],
       [9.46861640e-01, 5.31383598e-02],
       [8.91401040e-01, 1.08598960e-01],
       [8.46372308e-01, 1.53627692e-01],
       [9.21034863e-01, 7.89651375e-02],
       [9.41192245e-01, 5.88077548e-02],
       [9.35711095e-01, 6.42889049e-02],
       [9.65349997e-01, 3.46500032e-02],
       [9.34013106e-01, 6.59868936e-02],
       [9.38334315e-01, 6.16656848e-02],
       [8.25637381e-01, 1.74362619e-01],
       [8.43723225e-01, 1.56276775e-01],
       [9.10683645e-01, 8.93163547e-02],
       [8.96863190e-01, 1.03136810e-01],
```

```
[8.57051100e-01, 1.42948900e-01],  
[8.61150640e-01, 1.38849360e-01],  
[9.61546009e-01, 3.84539910e-02],  
[9.17038174e-01, 8.29618262e-02],  
[8.99207055e-01, 1.00792945e-01],  
[9.80184073e-01, 1.98159265e-02],  
[9.84855251e-01, 1.51447488e-02],  
[9.51013173e-01, 4.89868271e-02],  
[9.49128956e-01, 5.08710441e-02],  
[8.17197296e-01, 1.82802704e-01],  
[9.34439821e-01, 6.55601795e-02],  
[9.32979522e-01, 6.70204783e-02],  
[8.75432402e-01, 1.24567598e-01],  
[8.86008026e-01, 1.13991974e-01],  
[9.17893951e-01, 8.21060489e-02],  
[8.62235968e-01, 1.37764032e-01],  
[9.68257237e-01, 3.17427635e-02],  
[9.99355235e-01, 6.44764749e-04],  
[9.43561460e-01, 5.64385404e-02],  
[9.57346750e-01, 4.26532503e-02],  
[9.35292082e-01, 6.47079177e-02],  
[8.47691625e-01, 1.52308375e-01],  
[9.95317997e-01, 4.68200265e-03],  
[8.30706734e-01, 1.69293266e-01],  
[9.69221505e-01, 3.07784955e-02],  
[9.32373291e-01, 6.76267090e-02]])
```

```
plt.figure(1)  
plt.plot([0,1],[0,1],"k--")  
plt.plot(fpr_lr,tpr_lr,label = "LR")  
plt.xlabel("False positive rate")  
plt.ylabel("True positive rate")  
plt.title("ROC curve - LR Model")  
plt.legend(loc="best")  
plt.show()
```

ROC curve - LR Model

```

metrics.auc(fpr_lr,tpr_lr)

0.8125

forest = RandomForestClassifier()
forest.fit(X_train_processed, y_train)

importances = forest.feature_importances_

indices = np.argsort(importances)[::-1]

print("Feature importance ranking by Random Forest Model:")
for ind in range(X.shape[1]):
    print ("{0} : {1}".format(X_train_processed.columns[indices[ind]],round(importances[indices[ind]])))

Feature importance ranking by Random Forest Model:
is_canceled : 0.5446
lead_time : 0.0924
agent : 0.0475
arrival_date_week_number : 0.0375
stays_in_weekend_nights : 0.0329
booking_changes : 0.0322
arrival_date_day_of_month : 0.0316
adr : 0.0275
adults : 0.027
stays_in_week_nights : 0.0217
arrival_date_month_April : 0.0181
company : 0.0159
meal_encoded : 0.0149
arrival_date_month_June : 0.0103
is_repeated_guest : 0.0082
arrival_date_month_May : 0.0076
arrival_date_year : 0.0064
arrival_date_month_November : 0.0056
required_car_parking_spaces : 0.0049
total_of_special_requests : 0.0033
arrival_date_month_January : 0.0029
previous_bookings_not_canceled : 0.0025

import matplotlib.pyplot as plt
top_features = 10
feature_names = X_train_processed.columns[indices][:top_features]
importances = importances[:top_features]

plt.figure(figsize=(10, 6))
plt.bar(range(len(importances)), importances, color='skyblue', align='center')
plt.xticks(range(len(importances)), feature_names, rotation=45)
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.title('Top 10 Feature Importance by Random Forest Model')
plt.tight_layout()

```